

# Welcome to HW2

In this assignment you will be implementing a neural network in order to perform regression on the Airfoil Self-Noise data set. Remember to restart and run all cells before submission. Points will be deducted if you do not do this. When you are ready to submit, you can convert your notebook to a PDF file by printing the page either with `ctrl + p` or `command + p` and then saving as p1.pdf.

## The imports and helper functions should not be modified in any way.

```
In [86]: import torch
from torch import nn
from torch.utils.data import DataLoader
import numpy as np
from tqdm import tqdm
from matplotlib import pyplot as plt
```

```
In [87]: def evaluate(model, test_data):
    ...
    Do not modify this code.
    ...
    test_loader = DataLoader(Dataset(test_data), batch_size=1)
    loss_fn = torch.nn.MSELoss()
    with torch.no_grad():
        total_loss = 0
        for x, y in test_loader:
            pred = model(x)
            total_loss += loss_fn(pred, y.view(-1,1)).item()
    print("TOTAL EVALUATION LOSS: {:.5f}".format(total_loss))
```

```
In [88]: def plot_training_curves(train_loss, val_loss, loss_fn_name, reduction):
    ...
    Do not modify this code.
    ...
    fig, ax = plt.subplots(figsize=(8,6))
    ax.plot(train_loss, label="Train Loss")
    ax.plot(val_loss, label="Validation Loss")
    ax.legend(loc='best')
    ax.set_title("Loss During Training", fontsize=16)
    ax.set_xlabel("Epochs", fontsize=14)
    ax.set_ylabel("Loss: {}(reduction={})".format(loss_fn_name, reduction), fontsize=14)
    plt.savefig("./example_loss.pdf")
    plt.show()
```

## a) Implement your dataset object.

Do not modify the function definitions. Please note that the first five columns of the airfoil data are features and the last column is the target. Your dataset should have one attribute for the

features, one attribute for the targets, and should return the specified features and target in `__getitem__()` as separate values.

```
In [89]: class Dataset(torch.utils.data.Dataset):
    """Create your dataset here."""

    def __init__(self, airfoil_data):
        """
        Initialize your Dataset object with features and labels
        """
        ### Define your features and labels here
        self.features = airfoil_data[:,0:5]
        self.labels = airfoil_data[:,5]

    def __len__(self):
        ### Define the length of your data set
        return (len(self.features))

    def __getitem__(self, idx):
        ### Return the features and labels of your data for a given index
        feature = self.features[idx,0:5]
        label = self.labels[idx]

        feature = torch.tensor(feature, dtype=torch.float32)
        label = torch.tensor(label, dtype = torch.float32)

        return feature, label
```

## b) Implement the model architecture and forward function.

Do not modify the function definitions. You will need to define input, hidden, and output layers, as well as the activation function.

```
In [95]: class NeuralNetwork(nn.Module):
    def __init__(self, input_dimension=5, output_dimension=1, hidden=[32], activation=
        super(NeuralNetwork, self).__init__()
        ...
        Implement your neural network here. You will need to add layers and an act
        You are free to modify the number of layers in the hidden list.
        ...
        ### Define your input, hidden and output layers here
        self.lin1 = nn.Linear(input_dimension,hidden[0])
        self.lin2 = nn.Linear(hidden[0],hidden[1])
        self.lin3 = nn.Linear(hidden[1],hidden[2])
        self.lin4 = nn.Linear(hidden[2],hidden[3])
        self.lin5 = nn.Linear(hidden[3],hidden[4])
        self.lin6 = nn.Linear(hidden[4],output_dimension)
        self.act = activation
        ### Set your activation function here

    def forward(self, x):
        ...
```

```

    ... Implement the forward function using the layers and activation function you
    ...
    ### Call your hidden layers and activation function to do the forward pass through the model
    x = self.lin1(x)
    x = self.act(x)
    x = self.lin2(x)
    x = self.act(x)
    x = self.lin3(x)
    x = self.act(x)
    x = self.lin4(x)
    x = self.act(x)
    x = self.lin5(x)
    x = self.act(x)
    x = self.lin6(x)

    return x

```

## c, d) Define hyperparameters and implement the training loop.

You will need to choose your loss function, number of epochs, optimizer learning rate, optimizer weight decay, and batch size for part (c). You will need to set up the DataLoader, implement the forward pass, and implement the backpropagation update.

```

In [96]: def train(model, train_data, validation_data):

    ###
    # Modify these parameters
    ###
    loss_fn = nn.MSELoss()
    epochs = 3000
    learning_rate = 0.001
    weight_decay = 0.0001
    batch_size = 64

    # Set up data
    train_loader = DataLoader(Dataset(train_data), batch_size=batch_size)
    validation_loader = DataLoader(Dataset(validation_data), batch_size=batch_size)

    # The Adam optimizer is recommended for this assignment.
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)

    train_losses, val_losses = [], []
    for ep in tqdm(range(epochs)):
        train_loss = 0
        for x, y in train_loader:

            # Make prediction with your model

            pred = model(x)

            # Calculate Loss
            loss = loss_fn(pred, y.view(-1, 1))

            train_loss += loss.item()

```

```

        # Backpropagate Loss through the network and update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    val_loss = 0
    with torch.no_grad():
        for x, y in validation_loader:

            # Make prediction with model.forward()
            pred = model.forward(x)

            # Calculate Loss
            loss = loss_fn(pred, y.view(-1, 1))

            val_loss += loss.item()

    # Feel free to modify how frequently training progress is printed
    if(ep%500 == 0):
        print("Train Loss: {0:.4f}\tValidation Loss: {1:.4f}".format(train_loss, val_loss))

    # Hold on to Losses for easy saving and plotting
    train_losses.append(train_loss)
    val_losses.append(val_loss)

# Save your losses as .numpy files
np.save("./train_losses.npy", train_losses)
np.save("./val_losses.npy", val_losses)

# Save the model as ./p1_model.pt
torch.save(model.state_dict(), "./p1_model.pt")
return model

```

## e) Load your data, then train and evaluate your model before plotting the training curves.

```

In [97]: if __name__ == '__main__':
        torch.manual_seed(137)

        # Load in the provided data
        train_data = np.load('train_data.npy')
        validation_data = np.load('validation_data.npy')
        test_data = np.load('test_data.npy')

        model = NeuralNetwork(input_dimension=5, output_dimension=1, hidden=[128, 256, 512, 256])

        model = train(model, train_data, validation_data)
        evaluate(model, test_data)

        # Load your training data and call the provided plot function. Loss function and r
        # required for the plotting function.

        train_loss = np.load('train_losses.npy')
        val_loss = np.load('val_losses.npy')
        plot_training_curves(train_loss, val_loss, loss_fn_name="MSELoss", reduction="mean")


```

```

0%|          | 2/3000 [00:00<06:10, 8.09it/s]
Train Loss: 11.3520      Validation Loss: 2.4397
17%|██       | 502/3000 [00:57<04:39, 8.93it/s]
Train Loss: 0.1922      Validation Loss: 0.2784
33%|████     | 1002/3000 [01:57<03:57, 8.41it/s]
Train Loss: 0.1221      Validation Loss: 0.2745
50%|██████   | 1503/3000 [02:57<02:47, 8.96it/s]
Train Loss: 0.1045      Validation Loss: 0.3147
67%|████████ | 2002/3000 [03:55<01:56, 8.54it/s]
Train Loss: 0.1047      Validation Loss: 0.2668
83%|█████████| 2502/3000 [04:55<00:56, 8.81it/s]
Train Loss: 0.0584      Validation Loss: 0.2430
100%|██████████| 3000/3000 [05:59<00:00, 8.34it/s]
TOTAL EVALUATION LOSS: 11.13950

```

Loss During Training



**f Run 4 different hyperparameter combinations and explain the differences in results**

```

In [93]: # Changing four different hyperparameters
# 4 parameters are Learning rate, epoch, weight decay and batch size

def train(model, train_data, validation_data, lr = 0.005, epoch = 1000, weight_d = 0.0001

    loss_fn = nn.MSELoss()
    epochs = epoch

```

```

learning_rate = lr
weight_decay = weight_d
batch_size = batch

# Set up data
train_loader = DataLoader(Dataset(train_data), batch_size=batch_size)
validation_loader = DataLoader(Dataset(validation_data), batch_size=batch_size)

# The Adam optimizer is recommended for this assignment.
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)

train_losses, val_losses = [], []
for ep in tqdm(range(epochs)):
    train_loss = 0
    for x, y in train_loader:

        # Make prediction with your model

        pred = model(x)

        # Calculate loss
        loss = loss_fn(pred, y.view(-1, 1))

        train_loss += loss.item()

        # Backpropagate loss through the network and update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    val_loss = 0
    with torch.no_grad():
        for x, y in validation_loader:

            # Make prediction with model.forward()
            pred = model.forward(x)

            # Calculate loss
            loss = loss_fn(pred, y.view(-1, 1))

            val_loss += loss.item()

    # Feel free to modify how frequently training progress is printed
    if(ep%200 == 0):
        print("Train Loss: {0:.4f}\tValidation Loss: {1:.4f}".format(train_loss, val_loss))

    # Hold on to losses for easy saving and plotting
    train_losses.append(train_loss)
    val_losses.append(val_loss)

# Save your losses as .numpy files
np.save("./train_losses.npy", train_losses)
np.save("./val_losses.npy", val_losses)

# Save the model as ./p1_model.pt
torch.save(model.state_dict(), "./p1_model.pt")
return model

```

```

#parameters are learning rate, epochs, weight decay and batch size
parameters = [{'learning rate':0.01,'epochs':1000,'weight decay':0.0001,'batch size':100},
               {'learning rate':0.001,'epochs':2000,'weight decay':0.0001,'batch size':100},
               {'learning rate':0.001,'epochs':1000,'weight decay':0.01,'batch size':64},
               {'learning rate':0.001,'epochs':1000,'weight decay':0.0001,'batch size':100}]

if __name__ == '__main__':
    torch.manual_seed(137)

    for param in parameters:

        # Load in the provided data
        train_data = np.load('train_data.npy')
        validation_data = np.load('validation_data.npy')
        test_data = np.load('test_data.npy')

        model = NeuralNetwork(input_dimension=5,output_dimension=1,hidden=[128,256,512])

        model = train(model,train_data,validation_data,lr = param['learning rate'],epochs=param['epochs'],
                      weight_decay=param['weight decay'],batch_size = param['batch size'])

        evaluate(model,test_data)

        # Load your training data and call the provided plot function. Loss function and plot function
        # required for the plotting function.


        train_loss = np.load('train_losses.npy')
        val_loss = np.load('val_losses.npy')
        plot_training_curves(train_loss,val_loss,loss_fn_name="MSELoss",reduction="mean")

```

```

0%|          | 0/1000 [00:00<?, ?it/s] 0%|          | 2/1000 [00:00<02:43, 6.11it/s]
Train Loss: 25.8555      Validation Loss: 1.9970
20%|██       | 202/1000 [00:25<01:37, 8.18it/s]
Train Loss: 0.2317      Validation Loss: 0.2766
40%|████     | 402/1000 [00:52<01:32, 6.45it/s]
Train Loss: 0.1580      Validation Loss: 0.2437
60%|██████   | 602/1000 [01:17<00:46, 8.49it/s]
Train Loss: 0.1331      Validation Loss: 0.2504
80%|████████ | 802/1000 [01:44<00:27, 7.14it/s]
Train Loss: 7.5122      Validation Loss: 3.1158
100%|██████████| 1000/1000 [02:17<00:00, 7.26it/s]
TOTAL EVALUATION LOSS: 276.44625

```




```


0%|          | 1/2000 [00:00<09:35, 3.47it/s]
Train Loss: 17.9520    Validation Loss: 3.7931
10%|█         | 201/2000 [00:57<09:55, 3.02it/s]
Train Loss: 0.5322    Validation Loss: 0.7939
20%|██        | 402/2000 [01:47<05:04, 5.25it/s]
Train Loss: 0.3218    Validation Loss: 0.6576
30%|███       | 602/2000 [02:27<04:39, 5.01it/s]
Train Loss: 0.2216    Validation Loss: 0.6459
40%|████      | 802/2000 [03:07<03:58, 5.02it/s]
Train Loss: 0.3273    Validation Loss: 0.7036
50%|█████     | 1002/2000 [03:47<03:28, 4.79it/s]
Train Loss: 0.1837    Validation Loss: 0.5928
60%|██████    | 1202/2000 [04:29<02:42, 4.92it/s]
Train Loss: 0.2298    Validation Loss: 0.6312
70%|███████   | 1401/2000 [05:08<01:56, 5.12it/s]
Train Loss: 0.1467    Validation Loss: 0.5821
80%|████████  | 1601/2000 [05:50<01:21, 4.90it/s]
Train Loss: 0.0895    Validation Loss: 0.5829
90%|█████████ | 1801/2000 [06:30<00:47, 4.18it/s]
Train Loss: 0.1935    Validation Loss: 0.5732
100%|██████████| 2000/2000 [07:11<00:00, 4.64it/s]
TOTAL EVALUATION LOSS: 12.78716

```

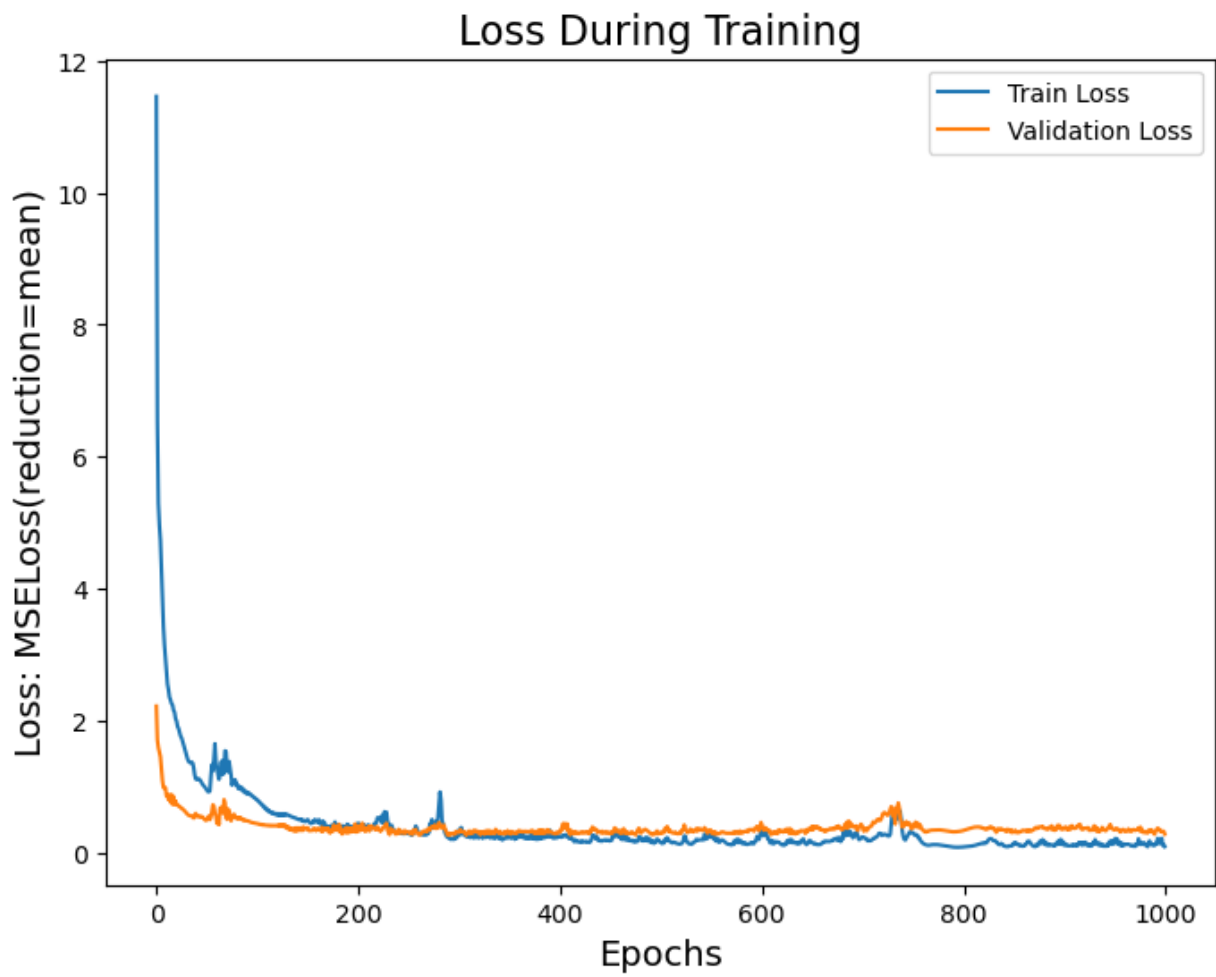




0%	2/1000	[00:00<02:33, 6.51it/s]
Train Loss: 13.9110      Validation Loss: 3.9300		
20%	202/1000	[00:26<02:01, 6.59it/s]
Train Loss: 2.4092      Validation Loss: 0.8889		
40%	402/1000	[00:53<01:16, 7.86it/s]
Train Loss: 2.1463      Validation Loss: 0.7920		
60%	602/1000	[01:21<01:00, 6.61it/s]
Train Loss: 2.0308      Validation Loss: 0.7582		
80%	802/1000	[01:48<00:25, 7.88it/s]
Train Loss: 1.9633      Validation Loss: 0.7430		
100%	1000/1000	[02:16<00:00, 7.34it/s]
TOTAL EVALUATION LOSS: 41.66925		



0%		2/1000 [00:00<02:19, 7.14it/s]
Train Loss: 11.4599      Validation Loss: 2.2203		
20%	█	202/1000 [00:25<01:38, 8.08it/s]
Train Loss: 0.4464      Validation Loss: 0.3269		
40%	██	402/1000 [00:50<01:16, 7.86it/s]
Train Loss: 0.2467      Validation Loss: 0.3098		
60%	████	602/1000 [01:16<00:50, 7.87it/s]
Train Loss: 0.3012      Validation Loss: 0.3880		
80%	██████	802/1000 [01:41<00:22, 8.85it/s]
Train Loss: 0.0883      Validation Loss: 0.3389		
100%	██████████	1000/1000 [02:06<00:00, 7.91it/s]
TOTAL EVALUATION LOSS: 11.80886		



## Explanation

The four parameters I made changes to are learning rate, epochs, weight decay and batch size. I used the same network and the loss function same for all the plots as MSELoss was giving the best results than any other losses.

For the first plot, I made the learning rate of 0.01 which is larger than the others which showed the total evaluation error of 276.45 which is too high. This is because the overshoot is happening and that leads to instability in convergence. The higher learning rate is leading to unstable training behavior therefore making it difficult for the model to predict how loss will change from one iteration to the next.

For the second plot, I increased the epochs to 2000 and reduced the batch size to 32 which gave me an error of 12.78 which is good but not the best. This is because the model to overfitting as the model trains for too many epochs without proper regularization. Along with that, due to smaller batch size result in noiser gradient estimates which makes the process less stable. This leads to more time in convergence.

For the third plot, I increased the weight decay to 0.01 and the batch size to 128 which gave me an error of 41. This can happen because high weight decay may have lead it to underfit where

the model could have become constrained and with higher batch size the generalization becomes difficult and it impacts the training.

For the fourth plot, all the values were small and this gave us an error of 11.8 which is pretty good as the model is trained properly but the plot has some disturbances because of which I felt maybe it is underfitting a little.

Based on all the assumptions the graph I made at first which is giving me the best model is the one with higher epoch of 3000 and a batch size of 64. The learning rate was 0.001 and the weight decay was 0.0001. This gave me the best model with a score of 11.1. Therefore, I chose these hyperparameters for this model.